

# Homework02

Yuchen Shao

Link to GitHub repository

<https://github.com/LAREINA-SHAO/STATS-506.git>

## Problem 1 - Dice Game

### a. Version 1

```
#' Title Dice Game Function - Using a loop
#
#' @param n The number of dice to roll
#
#' @return Total winnings
play_dice1 <- function(n){
  if (n < 0) {
    stop("Please input a non-negative number.")
  }

  # if no dice are rolled, return 0
  if (n == 0){
    return(0)
  }

  total_winnings <- 0
  # roll the dice n times
  die <- sample(1:6, n, replace = TRUE)

  for (i in 1:n) {
    # it costs $2 per game
    total_winnings <- total_winnings - 2
  }
}
```

```

    # on a roll of 3 or 5, win twice the roll
    if (die[i] == 3 | die[i] == 5) {
      total_winnings <- total_winnings + 2 * die[i]
    }
  }
  return(total_winnings)
}

```

### a. Version 2

```

#' Title Dice Game Function - using built-in R vectorized functions
#'
#' @param n The number of dice to roll
#'
#' @return Total winnings
play_dice2 <- function(n){
  if (n < 0) {
    stop("Please input a non-negative number.")
  }

  # if no dice are rolled, return 0
  if (n == 0){
    return(0)
  }

  # roll the dice n times
  die <- sample(1:6, n, replace = TRUE)

  #winnings for rolls of 3 or 5
  winnings <- ifelse(die == 3 | die == 5, 2 * die, 0)
  total_winnings <- + sum(winnings) - 2 * n
  return(total_winnings)
}

```

### a. Version 3

```

#' Title Dice Game Function - using table()
#'

```

```

#' @param n The number of dice to roll
#'
#' @return Total winnings
play_dice3 <- function(n){
  if (n < 0) {
    stop("Please input a non-negative number.")
  }

  # if no dice are rolled, return 0
  if (n == 0){
    return(0)
  }

  # roll the dice n times
  die <- sample(1:6, n, replace = TRUE)

  # create a frequency table (include 0 counts)
  die_table <- table(factor(die, levels = 1:6))

  # add together winnings and subtract out the total cost
  winnings_3 <- die_table["3"] * 2 * 3
  winnings_5 <- die_table["5"] * 2 * 5
  total_winnings <- winnings_3 + winnings_5 - 2 * n
  names(total_winnings) <- NULL
  return(total_winnings)
}

```

#### a. Version 4

```

#' Title Dice Game Function - using sapply
#'
#' @param n The number of dice to roll
#'
#' @return Total winnings
play_dice4 <- function(n){
  if (n < 0) {
    stop("Please input a non-negative number.")
  }

  # if no dice are rolled, return 0

```

```

if (n == 0){
  return(0)
}

# roll the dice n times
die <- sample(1:6, n, replace = TRUE)

# apply the game rules using sapply
winnings <- sapply(die, function(x) {
  if (x == 3 || x == 5) {
    return(2 * x)
  } else {
    return(0)
  }
})
total_winnings <- sum(winnings) - 2 * n
return(total_winnings)
}

```

**b.**

```

# total winnings for Version 1 with inputs of 3 and 3,000 dice rolls
c(play_dice1(3), play_dice1(3000))

```

```
[1] 0 2142
```

```

# total winnings for Version 2 with inputs of 3 and 3,000 dice rolls
c(play_dice2(3), play_dice2(3000))

```

```
[1] -6 2208
```

```

# total winnings for Version 3 with inputs of 3 and 3,000 dice rolls
c(play_dice3(3), play_dice3(3000))

```

```
[1] 10 1992
```

```
# total winnings for Version 4 with inputs of 3 and 3,000 dice rolls
c(play_dice4(3), play_dice4(3000))
```

```
[1] 20 2372
```

**c.**

```
# inputs 3
# set the seed for reproducibility
set.seed(111)
result1 <- play_dice1(3)

# reset the seed to ensure same random sequence
set.seed(111)
result2 <- play_dice2(3)

set.seed(111)
result3 <- play_dice3(3)

set.seed(111)
result4 <- play_dice4(3)

c(result1, result2, result3, result4)
```

```
[1] 0 0 0 0
```

```
# inputs 3000
set.seed(111)
result1 <- play_dice1(3000)

set.seed(111)
result2 <- play_dice2(3000)

set.seed(111)
result3 <- play_dice3(3000)

set.seed(111)
result4 <- play_dice4(3000)

c(result1, result2, result3, result4)
```

```
[1] 1918 1918 1918 1918
```

**d.**

```
#install.packages("microbenchmark")
library(microbenchmark)
```

```
#' Title Benchmarks for Dice Game Implementations
#
#' @param n The number of dice rolls
#
#' @return A `microbenchmark` object showing the performance of each version
BenchMarks <- function(n) {
  microbenchmark(
    play_dice1 = play_dice1(n),
    play_dice2 = play_dice2(n),
    play_dice3 = play_dice3(n),
    play_dice4 = play_dice4(n),
    times = 10
  )
}

set.seed(111)

print(BenchMarks(1000))
```

Unit: microseconds

	expr	min	lq	mean	median	uq	max	neval
play_dice1	267.383	270.512	273.3922	272.8575	275.514	283.957		10
play_dice2	106.325	118.091	148.4071	137.3160	163.523	256.714		10
play_dice3	199.137	210.658	238.6254	238.2575	266.285	276.775		10
play_dice4	720.776	729.446	797.7464	760.5480	820.630	977.404		10

```
print(BenchMarks(100000))
```

Unit: milliseconds

	expr	min	lq	mean	median	uq	max	neval
play_dice1	24.019454	26.136917	29.299269	28.420536	31.51776	39.91268		10
play_dice2	7.317755	7.956419	9.446681	9.725216	10.47302	11.81915		10

```
play_dice3  9.280632 10.741924 11.230799 11.306169 12.27452 12.44183 10
play_dice4 75.367440 89.587469 105.412463 97.680122 117.30530 171.67006 10
```


The function implemented using R's built-in vectorized functions is the fastest, while the function using `sapply()` performs the slowest. The function using a loop performs better than the one using `table()` when handling large inputs, but this advantage is negligible with smaller input sizes.

e.

```
#simulate 10000 times and pre-allocate a vector to store results
results <- numeric(10000)

for (i in 1:10000) {
  # play 1000 dice rolls in each simulation
  results[i] <- play_dice2(1000)
}

# Plot a histogram of the total winnings
hist(results, breaks = 50, main = "Monte Carlo Simulation of Total Winnings",
      xlab = "Total Winnings", col = "lightblue", border = "black")
abline(v = mean(results), col = "red", lwd = 2)
```



```
cat("Mean of total winnings:", mean(results), "\n")
```

Mean of total winnings: 665.2222

This is not a fair game. The mean of total winnings across 10,000 simulations is significantly above zero, indicating that the player, on average, makes a profit.

## Problem 2 - Linear Regression

a.

```
cars <- read.csv("cars.csv")
names(cars)
```

```
[1] "Dimensions.Height"
[2] "Dimensions.Length"
[3] "Dimensions.Width"
[4] "Engine.Information.Driveline"
[5] "Engine.Information.Engine.Type"
[6] "Engine.Information.Hybrid"
[7] "Engine.Information.Number.of.Forward.Gears"
[8] "Engine.Information.Transmission"
[9] "Fuel.Information.City.mpg"
[10] "Fuel.Information.Fuel.Type"
[11] "Fuel.Information.Highway.mpg"
[12] "Identification.Classification"
[13] "Identification.ID"
[14] "Identification.Make"
[15] "Identification.Model.Year"
[16] "Identification.Year"
[17] "Engine.Information.Engine.Statistics.Horsepower"
[18] "Engine.Information.Engine.Statistics.Torque"
```

```
# rename the columns
names(cars) <- c("height", "length", "width", "driveline", "engine_type",
                "hybrid", "forward_gears", "transmission", "city_mpg",
                "fuel_type", "highway_mpg", "classification", "ID", "make",
                "model_year", "year", "horsepower", "torque")
```

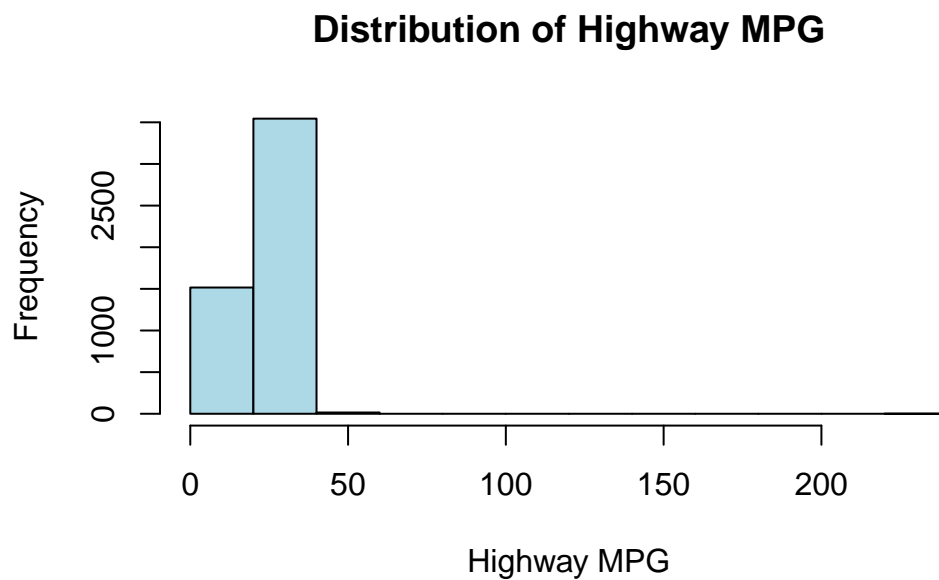


b.

```
gasoline_cars <- cars[cars$fuel == "Gasoline", ]
```

c.

```
# the distribution of highway_mpg  
hist(cars$highway_mpg, main = "Distribution of Highway MPG",  
      xlab = "Highway MPG", col = "lightblue", border = "black")
```

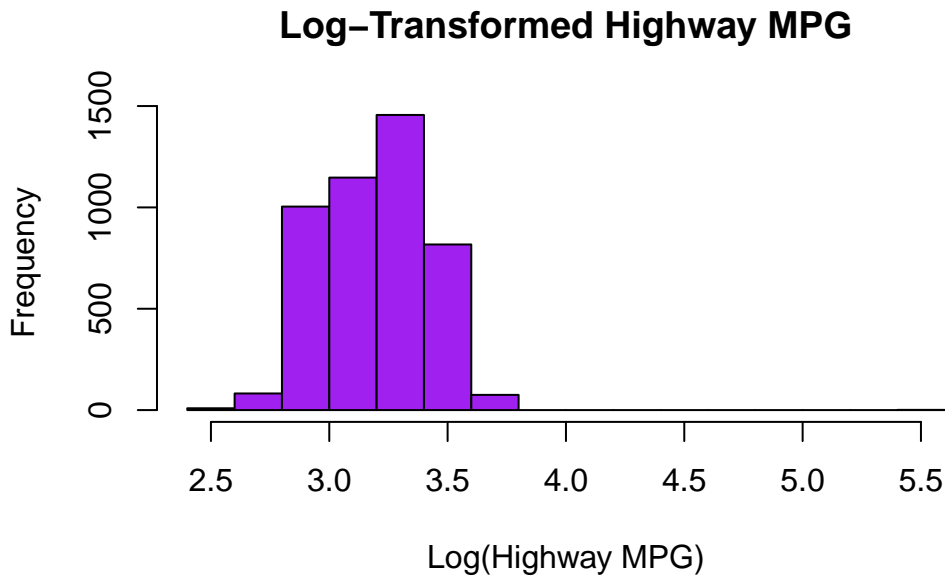


```
summary(gasoline_cars$highway_mpg)
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
13.00	21.00	25.00	24.97	28.00	223.00

```
# log transformation  
gasoline_cars$log_highway_mpg <- log(gasoline_cars$highway_mpg)
```

```
# the distribution of log_highway_mpg
hist(gasoline_cars$log_highway_mpg, main = "Log-Transformed Highway MPG",
     xlab = "Log(Highway MPG)", col = "purple", border = "black")
```



```
summary(gasoline_cars$log_highway_mpg)
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
2.565	3.045	3.219	3.194	3.332	5.407

The distribution of highway\_mpg is highly right-skewed, suggesting the presence of outliers and a long tail of higher values. To normalize the distribution, I applied a logarithmic transformation, which helped reduce skewness and made the data more suitable for analysis.

##d.

```
# convert 'year' to a factor (categorical variable)
gasoline_cars$year <- as.factor(gasoline_cars$year)
# fit a linear regression model
model <- lm(log_highway_mpg ~ torque + horsepower + height + length + width +
            year, data = gasoline_cars)
summary(model)
```

Call:

```
lm(formula = log_highway_mpg ~ torque + horsepower + height +  
    length + width + year, data = gasoline_cars)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-0.54759	-0.09385	-0.00414	0.09894	2.41852

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	3.507e+00	2.216e-02	158.236	< 2e-16 ***
torque	-2.294e-03	6.757e-05	-33.956	< 2e-16 ***
horsepower	9.238e-04	6.984e-05	13.227	< 2e-16 ***
height	4.050e-04	3.456e-05	11.719	< 2e-16 ***
length	3.475e-05	2.710e-05	1.282	0.19980
width	-8.722e-05	2.774e-05	-3.144	0.00168 **
year2010	-2.181e-02	2.076e-02	-1.051	0.29342
year2011	-2.430e-03	2.072e-02	-0.117	0.90665
year2012	4.012e-02	2.089e-02	1.921	0.05485 .

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.1412 on 4582 degrees of freedom

Multiple R-squared: 0.5638, Adjusted R-squared: 0.563


F-statistic: 740.3 on 8 and 4582 DF, p-value: < 2.2e-16

The estimated coefficient for torque is approximately -0.00229, meaning that for each additional unit of torque, the logarithm of highway MPG decreases by about 0.00229, holding all other variables constant. This indicates an inverse relationship between torque and highway fuel efficiency, where higher torque is associated with lower highway MPG. Additionally, this coefficient is highly significant, with a p-value less than 2e-16, confirming the robustness of the relationship.


e.

```
#install.packages("interactions")  
library(interactions)
```

```
hist(gasoline_cars$horsepower)
```



```
hist(gasoline_cars$torque)
```




```

model_interaction <- lm(log_highway_mpg ~ horsepower*torque + height + length +
                        width + as.factor(year), data = gasoline_cars)

interact_plot(model_interaction, pred = "torque", modx = "horsepower",
              # three different reasonable values of horsepower
              modx.values = c(150, 200, 250),
              at = list(year = 2010),
              main.title = "Interaction between Torque and Horsepower",
              xlab = "Torque",
              ylab = "Log Highway MPG",
              legend.main = "Horsepower Levels")

```

Using data `gasoline_cars` from global environment. This could cause incorrect results if `gasoline_cars` has been altered since the model was fit. You can manually provide the data to the `"data ="` argument.



The plot shows a negative relationship between torque and the logarithm of highway MPG, where increasing torque leads to a decrease in `log_highway_mpg`. The effect of torque on MPG is more pronounced for lower horsepower values, as shown by the steeper slopes for lower horsepower levels (150 and 200) compared to higher horsepower (250).

f.

```
# design matrix X for the linear regression
X <- model.matrix(log_highway_mpg ~ torque + horsepower + height + length +
                  width + as.factor(year), data = gasoline_cars)
y <- gasoline_cars$log_highway_mpg

# manually compute the OLS regression coefficients using matrix operations
# beta_hat = (X^T X)^(-1) X^T y
XtX_inv <- solve(t(X) %*% X)
XtY <- t(X) %*% y
beta_hat <- XtX_inv %*% XtY

cbind(model$coef, beta_hat)
```

	[,1]	[,2]
(Intercept)	3.506922e+00	3.506922e+00
torque	-2.294331e-03	-2.294331e-03
horsepower	9.238126e-04	9.238126e-04
height	4.049897e-04	4.049897e-04
length	3.475207e-05	3.475207e-05
width	-8.722295e-05	-8.722295e-05
year2010	-2.181247e-02	-2.181247e-02
year2011	-2.430359e-03	-2.430359e-03
year2012	4.011528e-02	4.011528e-02

The beta\_hat calculated manually is the same as lm did prior.